

Complexity of Linear Operators

Abstract

Let $A \in \{0, 1\}^{m \times n}$ be a matrix with t_0 zeroes and t_1 ones and \mathbf{x} be an n -dimensional vector over a semigroup. How many semigroup operations are required to compute $A\mathbf{x}$? This problem generalizes the well-known range queries problem and has applications in graph algorithms, functional programming languages, circuit complexity, and others. It is immediate that $O(t_1 + n + m)$ semigroup operations are sufficient. The main question studied in this paper is: can $A\mathbf{x}$ be computed using $O(t_0 + n + m)$ semiring operations? We prove that in general this is not possible: there exists a matrix $A \in \{0, 1\}^{n \times n}$ having exactly two zeroes in every row (hence $t_0 = 2n$) whose complexity is $\Theta(n\alpha(n))$. However, for the case when the underlying semiring is commutative, we prove an $O(t_0 + n + m)$ upper bound. This implies that for commutative setting, complements of sparse matrices can be processed as efficiently as sparse matrices (though the corresponding algorithm is more involved).

1 Introduction

1.1 Problem Statement and New Results

Let $A \in \{0, 1\}^{m \times n}$ be a matrix with t_0 zeroes and t_1 ones and $\mathbf{x} = (x_1, \dots, x_n)$ be an n -dimensional vector over a semigroup (S, \circ) . How many semigroup operations are required to compute the *linear operator* $A\mathbf{x}$? In this case, the i -th element of the output vector is $\sum_{j: A_{ij}=1} x_j$ where the summation is over the semigroup operation \circ . More generally, how many semigroup operations are needed to compute AB where $B \in S^{n \times k}$? In this paper, we are interested in lower and upper bounds involving t_0 and t_1 . Matrices with $t_1 = O(n)$ are usually called *dense* whereas matrices with $t_2 = O(n)$ are called *complements of dense matrices*. It is not difficult to see that computing all n outputs of $A\mathbf{x}$ independently takes $O(t_1 + n + m)$ semigroup operations. The main question studied in this paper is: can $A\mathbf{x}$ be computed using $O(t_0 + n + m)$ semiring operations? (This complexity is easy to achieve if \circ has an easily computable inverse (in that case S is a group): in this case A can be obtained by subtracting a dense matrix from all-ones matrix.) Our first result states that this is possible for *commutative* semigroups.

Theorem 1.1. *Let S be a commutative semigroup, $\mathbf{x} \in S^{n \times 1}$ be a vector, and $A \in \{0, 1\}^{m \times n}$ be a matrix with t_0 zeros. Then, $A\mathbf{x}$ can be computed using at most $O(m + n + t_0)$ semiring operations.*

As an immediate consequence we get the following matrix multiplication result.

Corollary 1.2. *Let S be a commutative semigroup, $A \in \{0, 1\}^{n \times n}$ be a matrix with t_0 zeros, and $B \in S^{n \times n}$ be a matrix. Then, AB can be computed using at most $O(n^2)$ semiring operations.*

We then show that commutativity is essential: for a strongly non-commutative semigroup S (the notion of strongly non-commutativity is made formal later in the text) the minimum number

of semigroup operations needed to compute $A\mathbf{x}$ for a matrix $A \in \{0, 1\}^{n \times n}$ with $t_0 = O(n)$ zeros is $\Theta(n\alpha(n))$ where α is the inverse Ackermann function.

Theorem 1.3. *For any strongly non-commutative semigroup X there is a circuit to compute any dense operator of size $O(n\alpha(n))$, where $\alpha(n)$ is the inverse Ackermann function. On the other hand, there exist dense matrices A such that any circuit computing Ax must have size $\Omega(n\alpha(n))$.*

1.2 Motivation

The linear operator problem is interesting for many reasons.

Range queries. In the *range queries* problem, one is given a vector $\mathbf{x} = (x_1, \dots, x_n)$ over a semiring (S, \circ) and a bunch of queries of the form (l, r) and is required to output the result $x_l \circ x_{l+1} \circ \dots \circ x_r$ for each query. The linear operator problem is thus a natural generalization of the range queries problem: each row of the matrix A defines a subset of the elements of \mathbf{x} that need to be summed up and this subset is not required to be a contiguous interval. The algorithms and hardness results for the linear operator problem presented in this paper are indeed inspired by some of the known results for the range queries problem. Later in the text we summarize a rich variety of algorithmic approaches and applications of the range queries problem.

Graph algorithms. Various graph path/reachability problems can be reduced naturally to matrix multiplication. Say, the all-pairs shortest path problem (APSP) is reducible to min-plus matrix product. Another example: the number of triangles in an undirected graph is equal to the trace of A^3 divided by six, where A is the adjacency matrix and matrix multiplication is over integers. It is natural to ask what happens if a graph has $O(n)$ edges or $O(n)$ anti-edges (as usual, by n we denote the number of nodes). In many cases, an efficient algorithm for sparse graphs is almost immediate whereas an algorithm with the same efficiency for a complement of a dense graph is not so immediate. For example, it is easy to solve APSP and triangle counting on dense graphs in time $O(n^2)$, but it is not immediate how to achieve the same time complexity for complements of dense graphs. In this paper, we show that the same effect occurs for linear operators: we prove that complements of dense matrices can be processed as efficiently as dense matrices though the corresponding algorithm is more involved. In particular, the efficient algorithm for matrix multiplication from Theorem ...

Matrix multiplication over semirings. Fast matrix multiplication methods rely essentially on the ring structure. The first such algorithm was given by Strassen [1], the current record upper bound is $O(n^{2.373})$ [2]; see the survey [3] for an overview of known approaches. For various important semirings lacking the inverse operation, we still do not know an $n^{3-\varepsilon}$ upper bound for a constant $\varepsilon > 0$. E.g., the strongest known upper bound for min-plus matrix multiplication is $n^3/\exp(\sqrt{\log n})$. In this paper, we present natural special cases of matrices where faster multiplication is possible.

Functional programming languages.

Circuit complexity. Computing linear operators over a Boolean semiring $(\{0, 1\}, \vee)$ is a well-studied problem in circuit complexity. The corresponding computational model is known as *rectifier network*. An overview of known lower and upper bounds for such circuits is given by Jukna [4, Section 13.6].

1.3 Organization of the Paper

2 (OLD) Introduction

Our main object of study in this paper are *dense matrices*, which we define as 0/1 matrices of size $n \times n$ that contain $O(n)$ zeroes. We are interested in computations over dense matrices that take place in algebraic structures lacking an inverse operation (e.g., semigroups and semirings). The interest in computations over such algebraic structures has recently grown substantially throughout the Computer Science community with the cases of Boolean and tropical semiring being of the main interest (see, for example, [?, ?, ?]). The lack of an inverse operation sometimes changes the complexity of algorithmic problems over algebraic structure drastically and even the complexity of standard computational tasks are not well understood over tropical and Boolean semirings (see, e.g. [?, ?]). From this perspective computations over dense matrices seems to be one of the most basic questions. In the presence of an inverse operation it trivially reduces to the computations over *sparse matrices*. Without an inverse operation dense matrices might intuitively seem harder than sparse matrices, but as we show in this paper, this intuition is only partially correct.

Consider a *dense linear operator*

$$\mathbf{y} = A\mathbf{x},$$

where A is a dense matrix and $\mathbf{x} = (x_1, \dots, x_n)$ is a vector over an arbitrary semigroup (see Section 3.1 for definitions and examples of algebraic structures used in this paper). (S, \circ) . Our goal is to simultaneously compute all elements of the resulting vector $\mathbf{y} = (y_1, \dots, y_n)$, where

$$y_i = \sum_{A_{ij}=1} x_j \tag{1}$$

for all $1 \leq i \leq n$, and the “summation” is over the semigroup operation \circ . What is the size of the smallest circuit comprising 2-input gates \circ that computes \mathbf{y} ?

Note that if A is sparse, i.e. contains $O(n)$ ones, then (1) directly yields a linear-size circuit. Furthermore, if the summation is over a *commutative group* rather than just a semigroup, then the dense case can be reduced to the sparse one by *subtracting* y_i from the sum $x_1 \circ \dots \circ x_n$. A similar reduction is also not hard to show for the case of non-commutative groups.

A natural solution that first comes to mind in the semigroup case is to split the rows of the matrix A into ranges of consecutive ones, thus obtaining $O(n)$ ranges overall, and then apply the classic *Range Queries* algorithm by Yao [?] to compute all ranges by a circuit of size $O(n\alpha(n))$, where $\alpha(n)$ is the inverse Ackermann function, and subsequently combine the results with $O(n)$ additional gates.

Can we do better? In general, the answer is “No”. However, if the semigroup is *commutative*, the answer is, remarkably, “Yes”! We present a linear-size circuit construction for dense linear operators in Section ???. Furthermore, in Section ?? we prove that the non-commutative case is equivalent to the Range Queries problem that, as shown by Chazelle and Rosenberg [?], requires $\Theta(n\alpha(n))$ gates, hence separating the complexity of the two cases.

Organization of the paper. Section 3 provides basic definitions used throughout the paper. Section 4 highlights applications of the presented linear-size construction. In particular, we observe that an $n \times n$ matrix can be multiplied by a dense matrix over an arbitrary semiring in $O(n^2)$ time. This appears to be significantly more challenging compared to the case of multiplication over a

ring, which can be reduced to sparse matrix multiplication via subtraction. In Section ?? we give the precise statements of our results.

3 Background

In this section we review basic algebraic structures used in this paper, as well as the classic Range Queries problem that turns out to be inherently related to dense linear operators.

3.1 Algebraic Structures

A *semigroup* (S, \circ) is an algebraic structure, where the operation \circ is *closed*, i.e., $\circ : S \times S \rightarrow S$, and *associative*, i.e., $x \circ (y \circ z) = (x \circ y) \circ z$ for all x, y , and z in S . *Commutative* (or *abelian*) semigroups introduce one extra requirement: $x \circ y = y \circ x$ for all x and y in S .

Commutative semigroups are ubiquitous. Below we list a few notable examples, starting with the most basic one, which is, arguably, known to every person on the planet.

- Integer numbers form commutative semigroups with many operations. For example, the order in which numbers are *added* is irrelevant, hence $(\mathbb{Z}, +)$ is a commutative semigroup. So are (\mathbb{Z}, \times) , (\mathbb{Z}, \min) and (\mathbb{Z}, \max) . On the other hand, it does matter in which order numbers are *subtracted*, hence $(\mathbb{Z}, -)$ is not a commutative semigroup: $1 - 2 \neq 2 - 1$. In fact, $(\mathbb{Z}, -)$ is not even a semigroup, since subtraction is non-associative: $1 - (2 - 3) \neq (1 - 2) - 3$.
- Boolean values form commutative semigroups (\mathbb{B}, \vee) , (\mathbb{B}, \wedge) , (\mathbb{B}, \oplus) and (\mathbb{B}, \equiv) .
- Any commutative semigroup (S, \circ) can be *lifted* to the set \hat{S} of “containers” of elements S , e.g., vectors or matrices, obtaining a commutative semigroup $(\hat{S}, \hat{\circ})$, where the lifted operation $\hat{\circ}$ is applied to the contents of containers element-wise. The lifting operation $\hat{\cdot}$ can often be omitted for clarity if there is no ambiguity.

The *average semigroup* $(\mathbb{Z} \times \mathbb{Z}, \circ)$ is a simple yet not entirely trivial example of semigroup lifting. By defining $(t_1, c_1) \circ (t_2, c_2) = (t_1 + t_2, c_1 + c_2)$, we can aggregate partial *totals* and *counts* of a set of numbers, which allows us to efficiently calculate their average as $avg(t, c) = \frac{t}{c}$. The average semigroup is commutative.

- Set union and intersection are commutative and associative operations giving rise to many set-based commutative semigroups. Here we highlight an example that motivated our research: the *graph overlay* operation, defined (this definition coincides with that of the *graph union* operation [?]). Graph union typically requires that given graphs are non-overlapping, hence it is not closed on the set of all graphs. Graph overlay does not have such a requirement, and is therefore closed and forms a semigroup). as $(V_1, E_1) + (V_2, E_2) = (V_1 \cup V_2, E_1 \cup E_2)$, where (V, E) is a standard set-based representation for directed unweighted graphs, comes from an algebra of graphs used in functional programming [?]. See further details in Section 4.3.

Groups extend semigroups by requiring the existence of the *identity element* $0 \in S$, such that $0 \circ x = x \circ 0 = x$, and the *inverse element* $-x \in S$ for all $x \in S$, such that $(-x) \circ x = x \circ (-x) = 0$. Groups provide a natural generalisation of arithmetic *subtraction*, whereby $x \circ (-y)$ denotes the subtraction of y from x .

A commutative semigroup (S, \circ) can often be extended to a *semiring* (S, \circ, \bullet) by introducing another associative (but not necessarily commutative) operation \bullet that *distributes* over \circ , that is

$$x \bullet (y \circ z) = (x \bullet y) \circ (x \bullet z)$$

$$(x \circ y) \bullet z = (x \bullet z) \circ (y \bullet z)$$

hold for all x, y , and z in S . Since \circ and \bullet behave similarly to numeric addition and multiplication, it is common to give \bullet a higher precedence to avoid unnecessary parentheses, and even omit \bullet from formulas altogether, replacing it by juxtaposition. This gives a terser and more convenient notation, e.g., the left distributivity law becomes: $x(y \circ z) = xy \circ xz$. We will use this notation, insofar as this does not lead to ambiguity.

Most definitions of semirings also require that the two operations have identities: the *additive identity*, denoted by 0, such that $0 \circ x = x \circ 0 = x$, and the *multiplicative identity*, denoted by 1, such that $1x = x1 = x$. Furthermore, 0 is typically required to be *annihilating*: $0x = x0 = 0$.

A semiring (S, \circ, \bullet) is also a *ring* if (S, \circ) is a group, i.e., the operation \circ is invertible. One can think of rings as semirings with subtraction.

Let us extend some of our semigroup examples to semirings:

- The most basic and widely known semiring is that of integer numbers with addition and multiplication: $(\mathbb{Z}, +, \times)$. Since every integer number $x \in \mathbb{Z}$ has an inverse $-x \in \mathbb{Z}$ with respect to the addition operation, $(\mathbb{Z}, +, \times)$ is also a ring. Interestingly, integer addition can also play the role of multiplication when combined with the max operation, resulting in the *tropical semiring* $(\mathbb{Z}, \max, +)$, which is also known as the *max-plus algebra*. Unlike $+$, the max operation has no inverse, therefore $(\mathbb{Z}, \max, +)$ is not a ring.
- Boolean values form the semiring $(\mathbb{B}, \vee, \wedge)$. Note that $(\mathbb{B}, \wedge, \vee)$ is a semiring too thanks to the duality between the operations \vee and \wedge . Furthermore, $(\mathbb{B}, \oplus, \wedge)$ is a ring, where every element is its own inverse: $x \oplus x = 0$ for $x \in \mathbb{B}$.
- Semirings and rings (S, \circ, \bullet) can also be lifted to the set \hat{S} of “containers” of elements S , most commonly matrices, obtaining $(\hat{S}, \hat{\circ}, \hat{\bullet})$. Matrices over tropical semirings, for example, are used for solving various path-finding problems on graphs.

3.2 Range Queries problem

Range Queries is a classical problem in data structures and algorithms having a variety of applications in fields like bioinformatics and string algorithms, computational geometry, image analysis, real-time systems, and others (we review some of the applications in Subsection ??, as well as a rich variety of algebraic techniques for the problem in Subsection ??).

In the Range Queries problem, one is given a sequence x_1, x_2, \dots, x_n of elements of a fixed semigroup (S, \circ) . Then, a *range query* is specified by a pair of indices (l, r) , such that $1 \leq l \leq r \leq n$. The answer to such a query is the result of applying the semigroup operation to the corresponding range, i.e., $x_l \circ x_{l+1} \circ \dots \circ x_r$. The Range Queries problem is then to simply answer all given range queries. There are two regimes: online and offline. In the *online regime*, one is given a sequence of *values* $x_1 = v_1, x_2 = v_2, \dots, x_n = v_n$ and is asked to preprocess it so that to answer efficiently any subsequent query. By “efficiently” one usually means in time independent of the length of the range (i.e., $r - l + 1$, the time of a naive algorithm), say, in time $O(\log n)$ or $O(1)$. In this paper,

we focus on the *offline* version, where one is given a sequence together with all the queries, and are interested in the minimum number of semigroup operations needed to answer all the queries. Moreover, we study a more general problem: we assume that x_1, \dots, x_n are formal variables rather than actual semigroup values. That is, we study the circuit size of the corresponding computational problem (the formal definition of the computational model is given later in the text).

4 Motivation and Applications

In this section we discuss our motivation and demonstrate two applications of the presented linear-size construction for a dense linear operator: fast multiplication of dense and *boring* matrices over arbitrary semirings (Section 4.2) and compact algebraic representation of dense graphs (Section 4.3).

4.1 Dense Operators

Throughout this section we consider $n \times n$ matrices over an arbitrary semiring (S, \circ, \bullet) , where the operations \circ and \bullet have identities 0 and 1, respectively.

A matrix is *sparse* if most of its elements are 0. To be more precise, we further assume that a sparse matrix has $O(n)$ non-zero elements. Sparse matrices arise in many applications, and can be multiplied by arbitrary vectors in $O(n)$ time and arbitrary matrices in $O(n^2)$ time (multiplication by an $n \times n$ matrix can be thought of as multiplication by n vectors).

A *0/1 matrix* is a matrix whose elements belong to the set $\{0, 1\}$. A 0/1 matrix is *dense* if it has $O(n)$ zero elements, i.e. most of its elements are 1.

Note that multiplication of dense matrices by vectors can be viewed as a special case of the Range Queries problem. Indeed, we can split the rows of a dense matrix into $O(n)$ ranges, compute answers to these range queries, and then recover the rows by combining the constituent ranges.

As mentioned above, computations on dense matrices over algebraic structures with inverse operations can often be reduced to computations on sparse matrices. However, the situation changes for computations over semigroups or semirings, which lack inverse operations. In such cases, the computation complexity of various matrix operations can differ significantly from more classical settings, which is a recurring topic in the recent years (see, for example, [?, ?, ?]). This paper provides further insight on this topic. As far as we know, the complexity of the problem under consideration was not known even for the simplest semigroups like (\mathbb{B}, \vee) .

4.2 Dense and boring matrix multiplication

Our first main result presented in Section ?? allows us to obtain a linear-size circuit for multiplying a 0/1 dense matrix of size $n \times n$ by a vector in an arbitrary semiring. Our construction is explicit and the corresponding algorithm takes $O(n^2)$ time (faster implementations are possible if the input matrix is provided in a compressed form). As a consequence, we can multiply a 0/1 dense matrix A by an arbitrary matrix B in $O(n^2)$ time as follows:

- Construct a linear-size circuit for the dense linear operator $A\mathbf{x}$. Time complexity: $O(n^2)$.
- Evaluate the circuit on all n columns of the matrix B . Each evaluation takes $O(n)$ time, hence the overall time complexity of this step is also $O(n^2)$.

Furthermore, by combining the algorithms for sparse and dense matrix multiplication, one can obtain an efficient algorithm for the multiplication of so-called *boring* matrices.

A matrix is *boring* if most of its elements are equal to some element b from the semiring. To be more precise, we further assume that a boring matrix has $O(n)$ elements that are not equal to b . Boring matrices are a natural generalisation of sparse and dense matrices: both are just special cases with $b = 0$ and $b = 1$, respectively.

To multiply a boring matrix A by a vector \mathbf{x} , we decompose the matrix into two components A_0 and A_1 , such that $A = A_0 \circ bA_1$, A_0 is sparse, and A_1 is dense (note that here the operations \circ and \bullet (the latter is represented by juxtaposition) are lifted to matrices). Now we can compute $A\mathbf{x}$ thanks to various semiring laws:

$$\begin{aligned} A\mathbf{x} &= (A_0 \circ bA_1)\mathbf{x} && \text{(sparse-dense decomposition)} \\ &= A_0\mathbf{x} \circ (bA_1)\mathbf{x} && \text{(distributivity and commutativity)} \\ &= A_0\mathbf{x} \circ b(A_1\mathbf{x}) && \text{(associativity)} \end{aligned}$$

Both $A_0\mathbf{x}$ and $A_1\mathbf{x}$ can be computed using sparse and dense matrix-vector multiplication, respectively; the results are further combined using scalar multiplication by b and vector addition \circ , both of which take linear time and have linear-size circuits. Note that the second step in the above equation relies on commutativity in a crucial way: elements of the original matrix A are partitioned into elements of A_0 and A_1 in an arbitrary order. As in the dense case, this immediately leads to $O(n^2)$ -time boring matrix multiplication.

4.3 Dense graph representation

Let us revisit the graph semigroup defined in Section 3.1. We will denote it by $(G_U, +)$, where G_U is the set of directed graphs whose vertices come from a universe U , that is, if $(V, E) \in G_U$ then $V \subseteq U$ and $E \subseteq V \times V$. Recall that the graph overlay operation $+$ is defined as

$$(V_1, E_1) + (V_2, E_2) = (V_1 \cup V_2, E_1 \cup E_2).$$

The algebra of graphs presented in [?] also defines the *graph connect* operation \rightarrow :

$$(V_1, E_1) \rightarrow (V_2, E_2) = (V_1 \cup V_2, E_1 \cup E_2 \cup V_1 \times V_2).$$

This operation allows us to “connect” two graphs by adding edges from every vertex in the left-hand graph to every vertex in the right-hand graph, possibly introducing self-loops if $V_1 \cap V_2 \neq \emptyset$. The operation is associative, non-commutative and distributes over $+$. Note, however, that the empty graph $\varepsilon = (\emptyset, \emptyset)$ is the identity for both overlay and connect operations: $\varepsilon + x = x + \varepsilon = x$ and $\varepsilon \rightarrow x = x \rightarrow \varepsilon = x$, and consequently the annihilating zero property does not hold, which makes this algebraic structure not a semiring according to the classic semiring definition.

By using the two operations one can construct any graph starting from primitive single-vertex graphs. For example, let $U = \{1, 2, 3\}$ and k stand for a single-vertex graph (k, \emptyset) . Then:

- $1 \rightarrow 2$ is the graph comprising a single edge $(1, 2)$, i.e. $1 \rightarrow 2 = (\{1, 2\}, \{(1, 2)\})$.
- $1 \rightarrow (2 + 3)$ is the graph $(\{1, 2, 3\}, \{(1, 2), (1, 3)\})$.
- $1 \rightarrow 2 \rightarrow 3$ is the graph $(\{1, 2, 3\}, \{(1, 2), (1, 3), (2, 3)\})$.

Clearly any sparse graph (V, E) , i.e. a graph with a sparse connectivity matrix, can be constructed by a linear-size expression:

$$(V, E) = \sum_{v \in V} v + \sum_{(u,v) \in E} u \rightarrow v.$$

But what about complements of sparse graphs, i.e. graphs with dense connectivity matrices? It is not difficult to show that by applying the dense linear operator we can obtain a linear-size circuit comprising 2-input gates $+$ and \rightarrow for any dense graph.

Let A be a dense matrix of size $n \times n$. Our goal is to construct the graph $G_A = (\{1, \dots, n\}, E)$ such that $(i, j) \in E$ whenever $A_{ij} = 1$.

First, we compute the dense linear operator $\mathbf{y} = A\mathbf{x}$ over the (commutative) graph semigroup $+$, where $\mathbf{x} = (1, 2, \dots, n)$, i.e., x_j is the primitive graph comprising a single vertex j , obtaining graphs y_i that comprise sets of isolated vertices corresponding to the rows of matrix A :

$$y_i = \sum_{A_{ij}=1} j.$$

The resulting graph $G_A = (\{1, \dots, n\}, E)$ can now be obtained by using the connect operation \rightarrow to connect i to all vertices y_i , and subsequently overlaying the results:

$$G_A = \sum_{i=1}^n i \rightarrow y_i.$$

Thanks to the linear-size construction for the dense linear operator, the size of the circuit computing G_A is $O(n)$. This allows us to store dense graphs on n vertices using $O(n)$ memory, and perform basic transformations of dense graphs in $O(n)$ time. We refer the reader to [?] for further details about applications of algebraic graphs in functional programming languages.

5 Computational Model

In this section we define our computational model, which includes a few less commonly known notions related to semigroups, as well as semigroup circuits.

5.1 Faithful semigroups

In the paper we consider computations over general semigroups. To establish complexity results and especially to establish reasonable lower bounds we need to consider semigroups with relatively rich structure.

Since we deal with computations with formal semigroup variables, it is convenient to introduce the following notation. Suppose (S, \circ) is a semigroup. Let $X_{S,n}$ be a semigroup with generators $\{x_1, \dots, x_n\}$ and with the equivalence relation consisting of identities in variables $\{x_1, \dots, x_n\}$ over (S, \circ) . That is, for two words W and W' in the alphabet $\{x_1, \dots, x_n\}$ we have $W = W'$ in $X_{S,n}$ iff no matter which elements of the semigroup S we substitute for $\{x_1, \dots, x_n\}$ we obtain a correct equation over S . In particular, note that if S is commutative (respectively, idempotent), then $X_{S,n}$ is also commutative (respectively, idempotent). We will often omit the subscript n and write simply X_S since the number of generators will be clear from the context.

Below we will use the following notation. Let W be a word in the alphabet $\{x_1, \dots, x_n\}$. Denote by $\text{Var}(W)$ the set of letters that are present in W .

The case of general commutative semigroups and computations over them was previously studied in relation to the Range Queries problem. The standard approach to capture their generality here is to consider *faithful commutative semigroups* [?, ?].

Definition 5.1. A commutative semigroup (S, \circ) is faithful commutative if for any equivalence $W \sim W'$ in X_S we have $\text{Var}(W) = \text{Var}(W')$.

Note that this definition does not pose any restrictions on the cardinality of each letter in W and W' . This allows us to capture in this definition important cases of idempotent semigroups. For example, semigroups $(\{0, 1\}, \vee)$ and (\mathbb{Z}, \min) are commutative faithful.

We also need to study the non-commutative case, and moreover, our results establish the difference between commutative and non-commutative cases. Thus, we need to extend the notion of faithfulness to non-commutative semigroups to capture their non-commutativity in the whole power. At the same time we would like to keep the case of idempotency. We introduce the notion of faithfulness for the non-commutative case inspired by the properties of free idempotent semigroups [?]. To introduce this notion we need several definitions.

The *initial mark* of W is the letter that is present in W such that its first appearance is farthest to the right. Let U be the prefix of W consisting of letters preceding the initial mark. That is, U is the maximal prefix of W with a smaller number of generators. We call U the *initial* of W . Analogously we define the *terminal mark* of W and the *terminal* of W .

Definition 5.2. We say that a semigroup X with generators $\{x_1, \dots, x_n\}$ is strongly non-commutative if for any words W and W' in the alphabet $\{x_1, \dots, x_n\}$ the equivalence $W \sim W'$ holds in X only if the initial marks of W and W' are the same, terminal marks are the same, the equivalence $U \sim U'$ holds in X , where U and U' are the initials of W and W' , respectively, and the equivalence $V \sim V'$ holds in X , where V and V' are the terminals of W and W' , respectively.

In other words, this definition states that the first and the last occurrences of generators in the equivalence separates the parts of the equivalence that cannot be affected by the rest of the generators and must therefore be the equivalences themselves. We also note that this definition exactly captures the idempotent case: for a free idempotent semigroup the condition in this definition is “if and only if” [?].

Definition 5.3. A semigroup (S, \circ) is faithful if X_S is strongly non-commutative.

We note that this notion of faithfulness is relatively general and is true for semigroups (S, \circ) with considerable degree of non-commutativity in their structure. It clearly captures free semigroups with at least two generators. It is also easy to see that the requirements in Definition 5.3 are satisfied for the free idempotent semigroup with n generators (if S is idempotent, then X_S is also clearly idempotent and no other relations are holding in X_S since we can substitute generators of S for x_1, \dots, x_n).

Next we observe some properties of strongly non-commutative semigroups that we need in our constructions.

Lemma 5.1. Suppose X is strongly non-commutative. Suppose the equivalence $W \sim W'$ holds in X and $|\text{Var}(W)| = |\text{Var}(W')| = k$. Suppose U and U' are minimal (maximal) prefixes of W and W' such that $|\text{Var}(U)| = |\text{Var}(U')| = l \leq k$. Then the equivalence $U \sim U'$ holds in X . The same is true for suffixes.

Proof. The proof is by induction on the decreasing l . Consider the maximal prefixes first. For $l = k$ and maximal prefixes we just have $U = W$ and $U' = W'$. Suppose the statement is true for some l , and denote the corresponding prefixes by U and U' , respectively. Then note that the maximal prefixes with $l - 1$ variables are initials of U and U' . And the statement follows by Definition 5.2.

The proof of the statement for minimal prefixes is completely analogous. Note that on the step of induction the prefixes differ from the previous case by one letter that are initial marks of the corresponding prefixes. So these additional letters are also equal by the Definition 5.2.

The case of suffixes is completely analogous. □

The next lemma is a simple corollary of Lemma 5.1.

Lemma 5.2. *Suppose X is strongly non-commutative. Suppose $W \sim W'$ holds in X_S . Let us write down the letters of W in the order in which they appear first time in W when we read it from left to right. Let's do the same for W' . Then we obtain exactly the same sequences of letters.*

The same is true if we read the words from right to left.

5.2 Circuits

We assume that the input consists of n formal variables $\{x_1, \dots, x_n\}$. We are interested in the minimum number of semigroup operations needed to compute all given words $\{w_1, \dots, w_m\}$ (e.g., for the range queries problem, each word has a form $x_l \circ x_{l+1} \circ \dots \circ x_r$). We use the following natural *circuit* model. A circuit computing all these queries is a directed acyclic graph. There are exactly n nodes of zero in-degree. They are labelled with $\{x_1, \dots, x_n\}$ and are called *input gates*. All other nodes have positive in-degree and are called *gates*. Finally, some m gates have out-degree 0 and are labelled as *output gates*. The *size* of a circuit is its number of edges (also called *wires*). Each gate of a circuit computes a word defined in a natural way: input gates compute just $\{x_1, \dots, x_n\}$; any other gate of in-degree r computes a word $f_1 \circ f_2 \circ \dots \circ f_r$ where $\{f_1, \dots, f_r\}$ are words computed at its predecessors (therefore, we assume that there is an underlying order on the incoming wires for each gate). We say that the circuit computes the words $\{w_1, \dots, w_m\}$ if the words computed at the output gates are equivalent to $\{w_1, \dots, w_m\}$.

For example, the following circuit computes range queries $(l_1, r_1) = (1, 4)$ and $(l_2, r_2) = (2, 5)$ over inputs $\{x_1, \dots, x_5\}$ or, equivalently, the linear operator $A\mathbf{x}$ where $A = \begin{pmatrix} 1 & 1 & 1 & 1 & 0 \\ 0 & 1 & 1 & 1 & 1 \end{pmatrix}$ and $\mathbf{x} = (x_1, \dots, x_5)^T$.

For any strongly non-commutative semigroup X there is a circuit to compute

For a 0/1-matrix A , by $C(A)$ we denote the *minimum number of gates* in a circuit computing the linear operator $A\mathbf{x}$.

A *binary circuit* is a circuit having no gates of fan-in more than two. It is not difficult to see that any circuit can be converted into a binary circuit of size at most twice the size of the original

circuit. For this, one just replaces every gate of fan-in k , for $k > 2$, by a binary tree with $2k - 2$ wires (such a tree contains k leaves hence $k - 1$ inner nodes and $2k - 2$ edges).

Clearly, in the binary circuit the number of gates does not exceed its size (i.e., the number of wires). And the number of gates in a binary circuit is exactly the minimum number of semigroup operations needed to compute the corresponding function.

Note that we can view circuits as computations over some semigroup (S, \circ) , meaning that we can substitute instead of the variables elements of the semigroup S . If we fix some semigroup (S, \circ) we can actually consider a circuit as a computation in the semigroups X_S . Moreover, we can forget about the original semigroup S and consider the computations in the circuit as computations in an arbitrary semigroup X with generators $\{x_1, \dots, x_n\}$.

In an important special case of the Boolean semigroup $(\{0, 1\}, \vee)$, circuits we are discussing are known as *rectifier networks*. An overview of known lower and upper bounds for such circuits is given by Jukna in [?, Section 13.6].